

Phage therapy for multidrug-resistant *Pseudomonas aeruginosa* infections: a systematic review of the structure and verifiability of the clinical evidence base

Danny Valdiviezo¹, Nataly Trelles¹

¹ School of Medicine, Universidad Católica de Cuenca, Cuenca, Ecuador.

Corresponding author: Danny Valdiviezo. Universidad Católica de Cuenca, Av. de las Américas y Humboldt, Cuenca, Ecuador. Email: dvchiqui[at]gmail.com

Word count: abstract 299; main text 4 284.

Tables: 4. **Figures:** 2. **Supplementary files:** 14 (S0–S13) plus a guide.

Abstract

Background. Carbapenem-resistant *Pseudomonas aeruginosa* is listed as a **high**-priority pathogen in the WHO 2024 list, having been critical priority in the 2017 list. Several recent syntheses have pooled phage therapy outcomes into overall success proportions. Doing so assumes this evidence base can be aggregated, and that assumption has never been examined.

Objectives. To delimit the clinical literature on phage therapy for resistant *P. aeruginosa* reproducibly, and to determine whether its structure and reporting completeness permit a quantitative synthesis of efficacy.

Methods. Systematic review following PRISMA 2020, with separate streams for databases and trial registers. Nine sources were searched without date restriction, including the regional databases BVS and SciELO. Only reports in English or Spanish were eligible, a criterion adopted after screening closed and verified report by report, including against the full text of every retrieved PDF. The unit of inclusion was the study, not the report. **Title and abstract screening was performed by a large language model acting as sole reviewer, without independent duplication;** data extraction is done by hand and in duplicate.

Results. From 23 057 records, 17 129 unique reports remained. Full-text assessment covered 233 reports, which grouped into 184 studies; 124 of these have a retrievable publication. Single case reports account for 39.5 % and only 18.5 % are comparative designs. Geographic origin is not stated in 61.3 % and resistance class cannot be assigned in 80.6 %. Full text was obtained for 91 studies (73.4 %); the remaining fraction concentrates 47.8 % of the comparative designs. The language criterion removed 35 studies, 18 of them comparative.

Conclusions. This evidence base is simultaneously large and structurally unsuited to a quantitative synthesis of efficacy. Published success proportions rest on assumptions these data do not support.

Keywords: bacteriophages; phage therapy; *Pseudomonas aeruginosa*; multidrug resistance; systematic review; PRISMA.

1. Introduction

Pseudomonas aeruginosa adds to its intrinsic resistance an unusual facility for acquiring new determinants. That combination keeps it among the World Health Organization's priority pathogens, though its position has moved: the 2017 list placed it in the critical tier ¹ and the 2024 update downgraded the carbapenem-resistant form to high priority ². The downgrade does not lessen the clinical problem an isolate with no options poses, but it does mean urgency should not be argued from a superseded classification. Once an isolate accumulates resistance to carbapenems, antipseudomonal cephalosporins, fluoroquinolones and aminoglycosides, the clinician's options can be counted on one hand. Sometimes none is both active and tolerable. The problem concentrates in device-associated infections, extensive burns, cystic fibrosis and bronchiectasis.

Phage therapy has re-emerged into that gap. Lytic bacteriophages act through a mechanism unrelated to antibiotic targets, spare the surrounding microbiota, and can be tailored to a patient's own isolate ³. But the clinical literature supporting them did not grow the way evidence for a drug in development grows. It grew as an accumulation of rescue cases published one at a time, almost always after the patient improved.

That growth pattern carries consequences that adding more cases does not dissolve. A published case report is, by construction, an event conditioned on its outcome. A compassionate-use series recruits patients in whom antibiotics have already failed and to whom the treating team judged it reasonable to offer something further⁴. And when phage is given alongside active antimicrobials, no favourable result can assign the credit to one or the other⁵.

Published syntheses have nonetheless tended to pool this literature into overall clinical success proportions^{6,7}. That step requires taking three things on trust that are rarely checked. First, that the resistance category used for stratification survives comparison with the criteria it claims to apply. Second, that the pooled outcome names the same thing across studies. Third, that the intervention is homogeneous enough for a common proportion to mean anything.

A prior question therefore precedes any estimate, and it is the one this review addresses: does this evidence base have the structure and reporting completeness that a quantitative synthesis requires? Answering it requires first delimiting the literature completely, including the regional literature and the trial registers that previous reviews have covered unevenly, and then characterising what that body of work states and what it leaves unsaid.

2. Methods

2.1 Protocol and registration

This review was conducted following the PRISMA 2020 statement⁸. The completed checklist accompanies the manuscript as supplementary material.

This review is not registered in PROSPERO or OSF. The absence is declared rather than omitted because prospective registration guards against *post hoc* modification of outcomes, and its absence is information readers need in order to calibrate the review. The protocol, the methodological decisions and their dates were recorded immutably in the working repository from the outset (section 2.8).

2.2 Eligibility criteria

Population. Human patients with *P. aeruginosa* infection classified by the authors as multidrug-resistant (MDR), extensively drug-resistant (XDR) or pandrug-resistant (PDR) according to Magiorakos et al.⁹, or described in terms implying those categories. Multi-pathogen studies were eligible when the *P. aeruginosa* subgroup was separable.

Intervention. Therapeutic administration of lytic bacteriophages by any route, alone or combined with antimicrobials. Endolysins and other phage-derived products given without the viral particle were excluded as a pharmacologically distinct intervention.

Comparator. None required, given the expected nature of the literature.

Outcomes. Clinical success, microbiological eradication, mortality, adverse events and emergence of phage resistance, as each study defined them.

Designs. Any design with treated patients: randomised and non-randomised trials, cohorts, case series and case reports. Excluded were laboratory, preclinical and modelling work; reviews without primary data of their own; secondary syntheses; veterinary infections; and studies of other organisms without a separable subgroup.

Language. Reports written in English or Spanish were eligible. **This criterion was not part of the original protocol and was adopted on 11 August 2026, with screening already closed**, so the decision was taken in knowledge of what the corpus contained. That is precisely what prospective registration exists to prevent, and it is therefore declared as an amendment, with its impact measured, in section 2.9 and among the limitations.

Language was not inferred from the journal name, which fails in both directions: *Pirogov Russian Journal of Surgery* publishes in English, and an English masthead can carry an article in another tongue. It was taken from the language field declared by Europe PMC, the same datum PubMed files under LA, and, where the source did not index it, from the language field of the export itself, which Scopus, SciELO and BVS all carry. Trial registry entries and Cochrane CENTRAL records were treated as eligible on the published

language rules of those sources rather than on assumption.

Verification did not stop at metadata. For the studies with a retrieved full text the language was checked against the **whole article**, analysing the central pages rather than the front matter, which many journals publish in English even when the body is not. That check exposed three reports whose metadata declared English and whose body is in Russian; all three were excluded. The 233 included reports rest on one of three classes of evidence: analysis of their own text, the language rule of the trial register that publishes them, or the source language field. None is unproven. The assignment for each report, with its evidence class and proof, is supplied as supplementary material.

Unit of inclusion. The **study**, not the report, following PRISMA 2020. A protocol, its conference abstracts and its final publication constitute one study with several reports. The decision and its justification are recorded in the project decision log.

2.3 Information sources and search strategy

Nine sources were searched across the two streams PRISMA 2020 requires to be kept separate. As **bibliographic databases**: PubMed, Scopus in two arms (one organism-first, one intervention-first), Cochrane CENTRAL, BVS including LILACS, CUMED, BINACIS and MedCarib, and SciELO. As **trial registers**: ClinicalTrials.gov, EudraCT and CTIS. No language, date or publication-type restriction was applied. The final run was on 10 August 2026.

The search strings were designed and run by D.V., an author of this review, querying each interface directly. They were neither generated nor executed by the language model whose involvement is declared in section 2.4: what to search for, in which sources and with what syntax was an author decision, taken before screening and fixed in writing. Full search strings are supplied as supplementary material, with the literal syntax of each interface, the number of hits and the date each was run, so that anyone can repeat them. Including BVS and SciELO was a deliberate choice. Clinical experience with phages published in Eastern Europe and Latin America is systematically under-represented in PubMed and Scopus, and leaving it out would have tilted the review towards the anglophone literature without saying so.

2.4 Study selection

Deduplication matched DOI, PMID and NCT transitively using union-find. Identifiers were read **only** from identifier fields, never from free text. The reason is concrete: in a trial run that scanned abstracts, four unrelated articles were merged because all four cited the same earlier paper. Records without an identifier were compared by normalised title and year, with no fuzzy matching.

Selection proceeded in three stages. **Stage 1** applied explicit, auditable rules: no title, publication type without primary data, animal scope declared in MeSH, and phage used as a laboratory tool. **Stage 2** screened titles and **stage 3** screened abstracts. Both used a **closed vocabulary of six exclusion codes** (ORG, REV, SEC, LAB, VET, OFF) fixed before screening began. A reason that fits none of the six does not receive an improvised seventh code: the record advances to the next stage instead of being excluded.

Who made each decision. Stage 1 is deterministic: a script executes it and every exclusion names the rule that fired, so anyone can re-run it and obtain the same result. Stages 2 and 3, the ones that require judgement, **were performed by a large language model — Claude Opus 5, Anthropic — acting as sole reviewer, without independent duplication**: 13 917 title decisions between 4 and 10 August 2026 and 460 abstract decisions on 5 August. The model received the title, abstract and metadata of each record, and could only choose among the six codes of the closed vocabulary. Every decision was logged with its reason, the model that issued it and its timestamp, in an append-only log supplied in full (S3).

No human reviewer adjudicated individual screening decisions, and agreement between the model and a human reader was not measured. That check is what duplicate screening provides, and it does not exist here. What does exist, and is what supports the claim that screening missed no eligible studies, is the positive-control audit of section 2.5: after each stage it was verified that none of the 40 studies known in advance lost all of its reports. A positive-control audit bounds false negatives over known studies; it does not bound them over unknown ones, and says nothing about false positives. Readers should calibrate the selection with that caveat, and section 4.4 records it as a limitation.

Data extraction, by contrast, is performed by the two authors in duplicate and independently, with no involvement of the model (section 2.6).

Decision logs are **append-only**. A correction is a new row superseding the previous one, and both remain. This review had to correct one erroneous exclusion, and that correction is auditable only because the original row survived.

2.5 Positive-control audit

Before screening began, a set of 40 studies known a priori to be eligible was fixed, drawn from an earlier scoping exercise. After each stage we verified that **no study in that set lost all of its reports**. The check operates at study level rather than report level, because excluding an erratum or a duplicate registry entry belonging to an included study is legitimate, whereas losing the study itself is not. The audit passed at all three stages (40 of 40). During one intermediate run it detected a miscalibrated exclusion rule, which was corrected before proceeding.

2.6 Data extraction

The extraction form, with 31 variables and controlled vocabularies, was piloted on five studies before use. Extraction is planned to be **independent and in duplicate** between the two reviewers, with disagreements resolved by consensus and agreement quantified variable by variable (Cohen's kappa).

Status at the time of this report. Duplicate extraction is under way and **has not concluded**. The data reported here come from a systematic pre-extraction from the abstracts of the 124 studies with a retrievable publication, performed with the same variable schema and explicitly flagged as partial in every record. Accordingly, **this report presents no efficacy estimates**. It presents the structure of the evidence base and its reporting completeness, which is what the pre-extraction does support. Outcome variables will be reported once duplicate extraction is complete.

2.7 Risk of bias and certainty

Risk of bias was to be assessed with RoB 2 for randomised trials¹⁰, ROBINS-I for non-randomised studies¹¹ and the Murad domains for case reports and series^{12,13}, and certainty with GRADE¹⁴. These assessments depend on full text and will be reported with the definitive extraction.

2.8 Synthesis methods and reproducibility

The feasibility of a meta-analysis of proportions was assessed explicitly against the usual requirements of that design^{15,16}: interpretable denominators, an outcome comparable across studies, and a sufficiently homogeneous intervention. The result of that assessment is presented in section 3.5.

The entire pipeline is written as versioned, re-runnable code: search, deduplication, screening, grouping and computation of every figure quoted. Each record is keyed on an identifier derived from its content and never on its position in a list, because a positional key changes referent as soon as the set of sources changes. An automated check reconciles the output of each stage against the input of the next, and fails if any declared source is not represented in the corpus. During development that check uncovered an entire database that had been searched and downloaded but never reached the corpus.

2.9 Protocol amendment: the language criterion

On 11 August 2026, with screening complete and 219 studies already delimited, eligibility was restricted to reports written in English or Spanish. The amendment answers an operational limit of the review team, which cannot extract reliably from a Russian text, and is declared here rather than presented as though it had applied from the outset.

Its effect was measured before it was adopted, and it was not small. It excluded **35 reports** and with them **35 whole studies**, none of which retained a report in an admissible language. The corpus fell from 219 to 184 studies and from 159 to 124 studies with a retrievable publication. The loss was not evenly spread: of 41 comparative designs, 23 remained, and of 21 randomised trials, 16.

Three of those 35 were detected only by reading the article itself. Their metadata declared English; the body is in Russian. The third surfaced only on the publisher's own page, which labels the full text as Russian while the abstract is in English. Had verification stopped at the source language field, all three would have stayed in the review.

The positive-control audit also recorded a concrete casualty. **Ronit et al. (2024)**, a case of phage therapy for a *P. aeruginosa*-infected vascular prosthesis published in Danish in *Ugeskrift for Læger*, belonged to the set of 40 studies fixed a priori as eligible and is excluded by the criterion. This is not a screening failure, since the study was identified and classified correctly, but the price of the criterion, and it is reported because an audit that is only reported when it passes demonstrates nothing.

The figures for the review without the amendment survive intact in the decision log, which is append-only, so any reader can reconstruct the earlier corpus and verify the effect independently.

3. Results

3.1 Study selection

The nine sources yielded **23 057 records**, from which 5 928 duplicates were removed to leave **17 129 unique reports** (Figure 1). Stage 1 rules excluded 3 235; of the 13 894 titles screened, 13 434 were excluded and 460 advanced to abstract screening; of those, 227 were excluded and **233 reports proceeded to full text**. Grouped by study, those 233 reports correspond to **184 studies**: 154 with a single report and 30 with several, up to a maximum of eleven reports of one trial.

Of the 184 studies, **124 have a retrievable publication**: 121 articles and 3 that exist only as conference abstracts. The remaining **60 are trial registry entries alone**, with no locatable associated publication. That figure by itself describes a field in which almost one study in three has been registered but has never published results.

Exclusion reasons at title screening were laboratory or preclinical work (4 062), other organism without a separable subgroup (3 031), review or commentary without primary data (2 807), does not evaluate phage therapy in patients (2 634), veterinary scope (838) and secondary synthesis (84). At abstract screening: other organism (65), review or commentary (64), does not evaluate phage therapy (40), **inadmissible language (35)**, laboratory (18) and secondary synthesis (5).

3.2 Full-text retrieval and its bias

Full text was obtained for **91 of the 124 retrievable studies (73.4 %)** through legitimate open-access routes. The remaining 33 require interlibrary loan, institutional access or a request to the authors.

The unretrieved fraction **is not a random sample of the whole** (Table 3). It contains 11 of the 23 comparative studies, **47.8 %**, and, although it reports fewer patients overall, it concentrates the designs that can sustain a comparison. Of the field's four landmark randomised trials^{17,18,19,20}, two remain unread: the intravesical phage trial and TP-102. The PhagoBurn article was obtained. The Yale trial has no full journal article, only a conference abstract and results deposited in the register; what was obtained is the trial protocol, which documents the design but not the outcomes.

The asymmetry is itself a finding. Full-text accessibility covaries with region of publication and with study design, and it pushes in the same direction as the language restriction: both tilt the corpus towards the anglophone open-access literature. A review that declares neither transforms the question without saying so.

3.3 Characteristics of the evidence base

The 124 retrievable studies were published between **2016 and 2026**, and **83.1 %** since 2020. This is a young and rapidly expanding evidence base (Table 1).

Its composition by design is what a rescue literature produces, not what a synthesis of efficacy needs: **49 single case reports (39.5 %)**, 19 case series, 16 randomised trials, 13 prospective cohorts, 7 non-randomised trials and 4 retrospective cohorts. A further 16 do not state a recognisable design in the abstract. Comparative designs total **23 studies, 18.5 %** of the whole, and are the only ones able to sustain a claim of relative efficacy.

Where stated, geographic origin is spread across Russia (7 studies), Poland and Georgia (6 each), Germany (4), France and the United States (3 each), and India, Iran and Israel (2 each). It is not stated in 76 studies. The fall in Russian representation relative to the pre-amendment corpus, from 34 studies to 7, is a direct effect of the language criterion and not of the search.

3.4 What the evidence base does not state

Reporting completeness, measured on what abstracts make explicit, is lowest precisely in the variables on which any stratification depends (Table 2):

- **Resistance class** (MDR, XDR or PDR) cannot be assigned in **80.6 %** of studies: 70.2 % do not mention it at all and a further 10.5 % mention it in terms that do not allow classification. For stratifying by resistance, which is what published syntheses do, both cases are equally unusable.
- **Route of administration** is not stated in **57.3 %**.
- **Pathogen scope**, that is, whether the study covers only *P. aeruginosa* or is mixed with a separable subgroup, is not stated in **52.4 %**.
- Fulfilment of the **DTR** criterion (difficult-to-treat resistance) is not stated in **98.4 %**.

These figures describe reporting in the abstract, not necessarily in the full article, and the distinction matters because some of the information may sit in the body of the text. That said, the abstract is what databases index and what automated reviews consume. Three studies in four failing to state their resistance category there means that no MDR/XDR/PDR stratification built from abstracts can be correct.

3.5 Feasibility of quantitative synthesis

A meta-analysis of proportions requires interpretable denominators, an outcome that names the same thing across studies, and a sufficiently homogeneous intervention. The body of evidence delimited here fails all three requirements.

Denominators. Almost four studies in ten are single cases, in which the success proportion can only take the values 0 or 1 and its binomial variance is maximal. Pooling forty-nine denominators of one with a handful of larger series produces a weighted mean whose weights are set by series size rather than by evidential quality.

Outcome. Among the studies whose full text has been read, "clinical success" neither names the same thing nor measures on the same scale. At one end, composite long-term clinical cure: no implant loosening and walking without pain at 21 months ²¹. At the other, discharge criteria: patient discharged, de-cannulated and rehabilitated ⁴. In between, the label without a criterion: "good clinical outcome" declared by the authors as binary, without stating what had to be observed ²².

The problem is not only heterogeneity. **Two of the landmark randomised trials report no binary success proportion at all:** the primary endpoint of PhagoBurn is the *time* to a sustained reduction in bacterial burden ¹⁷, and that of BX004-A is a *continuous* variable, the mean log-scale CFU reduction at day 15 ²³. A meta-analysis of proportions cannot include them without binarising what their authors measured otherwise, and that conversion would be made by the synthesiser, not by the study.

These definitions come from the full texts already retrieved, not from the pre-extraction from abstracts: **no abstract in the corpus states a definition of clinical success**, which is itself a datum about how the field reports.

Intervention. The products range from fixed-composition commercial polyvalent preparations to cocktails tailored to a patient's own isolate, and in most cases they are co-administered with active antimicrobials. The heterogeneity is conceptual rather than statistical, and no random-effects model resolves it: averaging different interventions yields a number that corresponds to no real intervention.

Added to this is the retrieval bias documented in section 3.2, which as things stand would exclude two-thirds of the comparative studies from any synthesis. **For these reasons no pooled efficacy proportion is presented here**, and we argue that those published in previous reviews rest on assumptions these data do not verify.

4. Discussion

4.1 Principal findings

This review delimits the clinical evidence base for phage therapy in resistant *P. aeruginosa* more broadly than previous syntheses: nine sources, no language restriction, regional literature and trial registers included. The result is 184 studies against the dozens handled by published reviews, and would have been 219 without the language restriction adopted at the end. That increase does not speak well of the field's maturity. Most of it comes from case reports, from registry entries without publication, and from regional literature that searches confined to PubMed and Scopus do not reach.

The central result is negative and worth stating plainly: **the breadth of this evidence base does not translate into capacity for synthesis**. Almost four retrievable studies in ten are a single case. Fewer than one in five is comparative. In four out of five, the resistance category by which stratification is attempted cannot be assigned. And almost one identified study in three exists only as a registry entry, with no published results.

4.2 Comparison with previous literature

Recent reviews have reported overall clinical success proportions consistently above 80 %^{6,7}. Those results do not contradict ours; they explain them. A literature composed largely of cases published after a favourable outcome yields, when pooled, a high proportion by construction. What that proportion measures is the propensity to publish successes, not the efficacy of the intervention.

Our delimitation further suggests that those syntheses worked on a fraction of the field. The unrestricted search retrieved 34 Russian and 3 Ukrainian studies, and that is precisely where comparative designs concentrate. The language amendment obliges us to set them aside, so this review documents that body of literature and where it sits but cannot speak to its content. That is an explicit debt to the next review with translation capacity.

4.3 Implications

For clinical practice. No published success proportion for phage therapy in resistant *P. aeruginosa* should be used to estimate the probability of benefit for an individual patient. The available evidence supports the statements that phage therapy is being given, that it is tolerated, and that favourable outcomes have been described. It does not support an effect size.

For research. Three shortcomings could be remedied at no added cost and would transform the usefulness of the next reports. State the resistance category together with the explicit criteria used to assign it. Document that the administered phage had verified lytic activity against the patient's isolate. And define clinical success with an explicit time horizon. The first is the most urgent: without it, the population of the question is undefined.

For registries. That 60 of 184 studies are registry entries without a locatable publication is a first-order publication-bias problem, and no statistical method resolves it.

4.4 Strengths and limitations

Strengths. The search covered nine sources across two separate streams, with no language restriction. Selection used a closed vocabulary fixed in advance and was audited against a positive-control set defined before screening began, which allows us to state, rather than merely hope, that screening lost no known eligible study. The whole pipeline is re-runnable code with content-derived keys, and an automated check reconciles each link with the next.

Limitations. There are five, and all are material. **First**, the review is not prospectively registered. **Second**, duplicate extraction has not concluded, so this report presents no outcomes, no risk-of-bias assessment and no GRADE certainty. It is a characterisation of the evidence base, not a synthesis of its results. **Third**, full text was obtained for 73.4 % of the retrievable studies, and the missing fraction is enriched in comparative designs, so what is described here under-represents the comparative part of the field. **Fourth**, title and abstract screening was issued by a large language model as sole reviewer, without independent duplication and without any human reviewer adjudicating individual decisions (section 2.4). Agreement between the model and a human reader was not measured, so it cannot be quantified how many eligible studies were missed or how many ineligible ones advanced. The positive-control audit bounds the former only over the 40 studies known in advance. It is the most serious limitation of this review and it affects how the corpus was delimited, not merely how it is described; data extraction is done in duplicate and by hand. **Fifth**, the restriction to English and Spanish was adopted once screening had closed, and removed 35 studies, 18 of them comparative and one belonging to the positive-control set. It pushes the corpus in the same direction as the retrieval bias and narrows what this review can say about the literature of Eastern Europe, which is where comparative designs concentrate.

5. Conclusions

The clinical evidence base for phage therapy in resistant *P. aeruginosa* is larger than published syntheses reflect and, at the same time, less suited to quantitative synthesis than those syntheses assume. It consists mainly of single cases and non-comparative series, published conditionally on their outcome, with a reporting completeness that does not allow verification of the resistance category defining the population of interest. While those three conditions hold, pooling their outcomes into a common proportion yields a figure that admits no clinical interpretation.

Declarations

Funding. This review received no specific funding from public, commercial or not-for-profit agencies.

Ethics approval and consent. Not applicable. This review synthesises already-published literature and aggregate data from public trial registers; it involves no patients, no specimens and no identifiable individual data, and therefore required neither ethics committee approval nor informed consent. No formal waiver was sought because the institution does not require one for studies without primary data.

Competing interests. The authors declare no competing interests.

Author contributions. [To be completed under CRediT before submission.]

Data and code availability. All material supporting this review is deposited in a public, citable repository: the search corpus with the raw exports of all nine sources, the append-only decision logs of the three screening stages, both reviewers' extraction workbooks, the methodological decision records with their dates, and the complete pipeline code, which is re-executable from search to the figures in the manuscript. «DEPOSIT DOI — TO BE COMPLETED BEFORE SUBMISSION». The decision logs are append-only and preserve every correction alongside the row it supersedes.

Use of artificial intelligence. A generative language model (Claude Opus 5, Anthropic) was used in four tasks, and they must be distinguished because they do not commit the work equally. *Issuing judgements that determine the corpus:* the model screened titles and abstracts as **sole reviewer, without independent duplication** --13 917 title decisions and 460 abstract decisions -- and performed the pre-extraction from abstracts. These decisions **were not individually adjudicated by a human reviewer**. Each is logged with its reason, the model and its timestamp, and the logs are supplied in full (S3, S4); the detail and its consequences are in sections 2.4 and 4.4. *Without issuing judgements:* assistance in programming the pipeline and in drafting. *Without any involvement of the model:* the design and execution of the searches, carried out by D.V. querying each interface (section 2.3); the eligibility criteria and the exclusion vocabulary, fixed before screening; and data extraction, performed by the two authors in duplicate. The authors designed the review, fixed the criteria and the exclusion vocabulary before starting, perform data extraction in duplicate without the model, and take full responsibility for the accuracy and integrity of the work. No AI tool is listed as an author, per ICMJE recommendations.

Acknowledgements. [To be completed.]

Tables and figures

The four tables are supplied as separate files (paper/tablas/), each with its own footnote. Both figures are supplied as vector PDF (paper/figuras/).

Table 1. Characteristics of the retrievable evidence base (n = 124 studies).

Table 2. Reporting completeness for variables critical to stratification.

Table 3. Comparison of studies with and without full text obtained.

Table 4. Exclusion reasons by stage, using the closed vocabulary.




Figure 1. PRISMA 2020 flow diagram.

What it shows. The passage from identified records to included studies, with the two identification streams — bibliographic databases and trial registers — counted separately, as the PRISMA 2020 statement requires. *How to read it.* The left column is the main flow; the grey boxes on the right hold what was excluded at each step, broken down by the closed six-code vocabulary (ORG, REV, SEC, LAB, VET, OFF; defined in Table 4). The register stream enters from the side and does not pass through bibliographic deduplication, because a registry entry has no DOI with which to collide against an article. The final grey box is not part of the selection flow: it reports how many included studies have been read in full text to date. *Source.* Project screening pipeline; final search run, 10 August 2026.




Figure 2. Composition of the retrievable evidence base (n = 124 studies).

What it shows. Panel A orders studies by design; Panel B distributes them by year of publication, separating comparative designs from the rest within each year. *How to read it.* In both panels the dark colour marks comparative designs — randomised and non-randomised trials — the only ones able to sustain a claim of relative efficacy. "Not stated" in Panel A means the abstract does not allow the design to be recognised, not that the study lacks one. The year 2026 is incomplete: the search closed in August. *Source.* Systematic pre-extraction from the abstracts of the 124 studies with a retrievable publication.

Supplementary material

- **S0.** Informe editorial: qué se decidió, cuándo y con qué fundamento.
- **S1.** Lista de comprobación PRISMA 2020.
- **S2.** Ecuaciones de búsqueda literales por fuente, con fecha y número de resultados.
- **S3.** Registro completo de decisiones de cribado, en dos ficheros: etapa 2 (título) y etapa 3 (resumen). Cada fila lleva su motivo codificado, quién la emitió y su marca de tiempo.
- **S4.** Pre-extracción desde el resumen, marcada como parcial en cada registro.
- **S5.** Listado de los 184 estudios con su situación y sus informes agrupados.
- **S6.** Auditoría de controles positivos.
- **S7.** Declaraciones ICMJE.
- **S8.** Registro de recuperación de texto completo: qué se intentó por cada estudio y con qué resultado.
- **S9.** Idioma por informe, con la clase de evidencia que aportaba.
- **S10.** Verificación del idioma sobre el texto completo del PDF, no sobre los metadatos.
- **S11.** Concordancia entre las dos extracciones independientes, antes de resolver los desacuerdos.
- **S12.** Resolución de conflictos de extracción, incluidos los que no pueden dirimirse todavía.
- **S13.** Reglas de extracción de desenlaces, con su corrección documentada.

Los acompaña una guía que explica, fichero a fichero, qué contiene y qué pregunta permite contestar.

References

1. Tacconelli E, Carrara E, Savoldi A, Harbarth S, Mendelson M, Monnet DL, et al. Discovery, Research, and Development of New Antibiotics: The WHO Priority List of Antibiotic-Resistant Bacteria and Tuberculosis. *Lancet Infectious Diseases*. 2018;18(3):318-327. doi:10.1016/S1473-3099(17)30753-3

2. World Health Organization. WHO Bacterial Priority Pathogens List, 2024: Bacterial Pathogens of Public Health Importance to Guide Research, Development and Strategies to Prevent and Control Antimicrobial Resistance. 2024.
3. Pirnay J, Djebara S, Steurs G, Griselain J, Cochez C, De Soir S, et al. Personalized Bacteriophage Therapy Outcomes for 100 Consecutive Cases: A Multicentre, Multinational, Retrospective Observational Study. *Nature Microbiology*. 2024;9(6):1434-1453. doi:10.1038/s41564-024-01705-x
4. Aslam S, Courtwright AM, Koval C, Lehman SM, Morales S, Furr C, et al. Early Clinical Experience of Bacteriophage Therapy in 3 Lung Transplant Recipients. *American Journal of Transplantation*. 2019;19(9):2631-2639. doi:10.1111/ajt.15503
5. Chan BK, Turner PE, Kim S, Mojibian HR, Eleftheriades JA, Narayan D. Phage Treatment of an Aortic Graft Infected With *Pseudomonas aeruginosa*. *Evolution, Medicine, and Public Health*. 2018;2018(1):60-66. doi:10.1093/emph/eoy005
6. Uyttebroek S, Chen B, Onsea J, Ruythooren F, Debaveye Y, Devolder D, et al. Safety and Efficacy of Phage Therapy in Difficult-to-Treat Infections: A Systematic Review. *Lancet Infectious Diseases*. 2022;22(8):e208-e220. doi:10.1016/S1473-3099(21)00612-5
7. Liu Y, Thong S, Moreira W, Yeo JH, Zhong Y, Chong ZS, et al. Clinical Application of Customized and Non-Customized Bacteriophage Therapy in Patients with Refractory/Resistant Bacterial Infections: A Systematic Review and Meta-Analysis. *International Journal of Antimicrobial Agents*. 2025;66(4):107570. doi:10.1016/j.ijantimicag.2025.107570
8. Page MJ, McKenzie JE, Bossuyt PM, Boutron I, Hoffmann TC, Mulrow CD, et al. The PRISMA 2020 Statement: An Updated Guideline for Reporting Systematic Reviews. *BMJ*. 2021;372:n71. doi:10.1136/bmj.n71
9. Magiorakos A, Srinivasan A, Carey RB, Carmeli Y, Falagas ME, Giske CG, et al. Multidrug-Resistant, Extensively Drug-Resistant and Pandrug-Resistant Bacteria: An International Expert Proposal for Interim Standard Definitions for Acquired Resistance. *Clinical Microbiology and Infection*. 2012;18(3):268-281. doi:10.1111/j.1469-0691.2011.03570.x
10. Sterne JAC, Savovi'c J, Page MJ, Elbers RG, Blencowe NS, Boutron I, et al. RoB 2: A Revised Tool for Assessing Risk of Bias in Randomised Trials. *BMJ*. 2019;366:l4898. doi:10.1136/bmj.l4898
11. Sterne JA, Hernan MA, Reeves BC, Savovi'c J, Berkman ND, Viswanathan M, et al. ROBINS-I: A Tool for Assessing Risk of Bias in Non-Randomised Studies of Interventions. *BMJ*. 2016;355:i4919. doi:10.1136/bmj.i4919
12. Murad MH, Sultan S, Haffar S, Bazerbachi F. Methodological Quality and Synthesis of Case Series and Case Reports. *BMJ Evidence-Based Medicine*. 2018;23(2):60-63. doi:10.1136/bmjebm-2017-110853
13. Munn Z, Barker TH, Moola S, Tufanaru C, Stern C, McArthur A, et al. Methodological Quality of Case Series Studies: An Introduction to the JBI Critical Appraisal Tool. *JBIC Evidence Synthesis*. 2020;18(10):2127-2133. doi:10.11124/JBISIR-D-19-00099
14. Guyatt G, Oxman AD, Akl EA, Kunz R, Vist G, Brozek J, et al. GRADE Guidelines: 1. Introduction --- GRADE Evidence Profiles and Summary of Findings Tables. *Journal of Clinical Epidemiology*. 2011;64(4):383-394. doi:10.1016/j.jclinepi.2010.04.026
15. Barker TH, Migliavaca CB, Stein C, Colpani V, Falavigna M, Aromataris E, et al. Conducting Proportional Meta-Analysis in Different Types of Systematic Reviews: A Guide for Synthesisers of Evidence. *BMC Medical Research Methodology*. 2021;21(1):189. doi:10.1186/s12874-021-01381-z
16. Migliavaca CB, Stein C, Colpani V, Barker TH, Ziegelmann PK, Munn Z, et al. Meta-Analysis of Prevalence: I^2 Statistic and How to Deal With Heterogeneity. *Research Synthesis Methods*. 2022;13(3):363-367. doi:10.1002/jrsm.1547
17. Jault P, Leclerc T, Jennes S, Pirnay J, Que Y, Resch G, et al. Efficacy and Tolerability of a Cocktail of Bacteriophages to Treat Burn Wounds Infected by *Pseudomonas aeruginosa* (PhagoBurn): A Randomised, Controlled, Double-Blind Phase 1/2 Trial. *Lancet Infectious Diseases*. 2019;19(1):35-45. doi:10.1016/S1473-3099(18)30482-1
18. Leitner L, Ujmajuridze A, Chanishvili N, Goderdzishvili M, Chkonia I, Rigvava S, et al. Intravesical Bacteriophages for Treating Urinary Tract Infections in Patients Undergoing Transurethral Resection of the Prostate: A Randomised, Placebo-Controlled, Double-Blind Clinical Trial. *Lancet Infectious Diseases*. 2021;21(3):427-436. doi:10.1016/S1473-3099(20)30330-3
19. Nir-Paz R, Onallah H, Dekel M, Gellman Y, Haze A, Ben-Ami R, et al. Randomized double-blind study on safety and tolerability of TP-102 phage cocktail in patients with infected and non-infected diabetic foot ulcers. *Med*. 2025;6:100565. doi:10.1016/j.medj.2024.11.018
20. Chan B, Kortright K, Stanley GL, Cochrane C, Lee A, Vill A, et al. CYstic Fibrosis bacterioPHage study at Yale (CYPHY). *Journal of Cystic Fibrosis*. 2023;22:S22. doi:10.1016/S1569-1993(23)00976-1

21. Ferry T, Kolenda C, Laurent F, Leboucher G, Merabishvili M, Djebara S, et al. Personalized Bacteriophage Therapy to Treat Pandrug-Resistant Spinal *Pseudomonas aeruginosa* Infection. *Nature Communications*. 2022;13:4239. doi:10.1038/s41467-022-31837-9
22. Onallah H, Hazan R, Nir-Paz R, others. Refractory *Pseudomonas aeruginosa* Infections Treated with Phage PASA16: A Compassionate Use Case Series. *Med*. 2023;4(9):600-611.e4. doi:10.1016/j.medj.2023.07.002
23. Weiner I, Kahan-Hanum M, Buchstab N, Zelcbuch L, Navok S, Sherman I, et al. Phage Therapy with Nebulized Cocktail BX004-A for Chronic *Pseudomonas aeruginosa* Infections in Cystic Fibrosis: A Randomized First-in-Human Trial. *Nature Communications*. 2025;16(1):5579. doi:10.1038/s41467-025-60598-4